UNIVERSITY OF PISA
Department of Computer Science
Master Programme in Data Science and Business Informatics

MASTER THESIS

**Minimizing Entropy for Training and Quantization (METaQ):
A novel algorithm for neural network training and compression**

SUPERVISORS
Prof. Antonio FRANGIONI
Prof. Paolo FERRAGINA

CANDIDATE
Andrea CARDIA

ACADEMIC YEAR 2023-24

Abstract

A neural network compression strategy is developed during the training phase by adding a regularization term $\phi(w)$ to the loss function, in order to minimize the entropy of the network weights. $\phi(w)$ results from a nondifferentiable optimization problem that is computationally very complex to solve. However, fortunately, to train the network, it is not necessary to explicitly find $\phi(w)$; instead, it is sufficient to provide its (sub)gradient to standard machine learning tools (such as PyTorch) to guide the training towards a low entropy weight configuration (in addition to achieving good accuracy). This subgradient can be computed using the optimal Lagrange multipliers β^* associated with the set of constraints involving the weights w in the problem that defines $\phi(w)$.

In this work, we will develop a procedure to determine β^* by applying Lagrangian relaxation and optimization techniques, also developing ad hoc methods for certain sub-problems when necessary for efficiency reasons. Once a trained network with low entropy is obtained, the compression strategy culminates in the quantization of the weights. The task of encoding the weights is implemented via well-known compression algorithms that come arbitrarily close to the entropy.

The introduction of the term $\phi(w)$ appropriately modeled for the function to be optimized during training is what makes this work innovative.

The tests were conducted using the LeNet-5 network on the MNIST dataset, although the strategy is also applicable to larger networks. The achieved results show a $29\times$ compression of LeNet-5 (Compression Ratio 3.43%) with an accuracy of 99.01%, making METaQ a strategy comparable to state-of-the-art NN-compression algorithms.

The project's source code is freely available at: https://github.com/cardiaa/METaQ

Contents

1 INTRODUCTION

Artificial Neural Networks (ANN) are nonlinear statistical models [Hastie et al., 2017] inspired by the functioning of the human brain, designed to process information in a way similar to how biological neurons process signals. These systems consist of a set of nodes, called artificial neurons, organized into layers and interconnected by edges, each associated with a weight.

Formally, denoting by a and b two neurons in the l-th layer, we can indicate the weight associated with their connection as $w_{a,b}^{(l)}$.[1]


Figure 1: Example of a neural network architecture with 3 inputs and 2 target classes. The set of connections between the input layer and the first hidden layer is represented by the matrix $w^{(1)}$. Source: mathworks.com

The set of these objects forms a tensor corresponding to the collection of weight matrices for each layer. Often, the indices a, b, and l are omitted unless strictly necessary, and the entire set of network weights is simply denoted as w to keep the notation concise. In the remainder of this work, we follow this convention, referring to n as the total number of weights in the network. Indicating with I the set of weights, we have $n = |I|$. We use the index i to refer to a certain weight w_i, underlying an unstructured view of the network.

The network receives input data in the first layer (called the input layer), where the data is appropriately organized into N records whose structure is to be captured. It then performs a dot product with the weights of the current layer, followed by a mathematical transformation, typically through an activation function, which is nonlinear to detect potential nonlinearity in the problem. The result is then transmitted to the neurons of the next layer. At the final layer, the output is the value $\hat{y}_j(w)$ corresponding to record j, with $j = 1, ..., N$, based on the given weight configuration $w = \{w_i\}_{i \in I} = \{w_i\}_{i=1}^n$.

The goal of training techniques is to minimize the error made by the network in its predictions compared to the true value of the function being modeled. This error is measured through a mathematical function, known as the loss function, which evaluates the discrepancy between the network's

---

[1]Without loss of generality, throughout the rest of the discussion, we will use the term `weight` to also refer to the concept of `bias`.

output and the expected result. This function, denoted as $L(w)$, depends on the weight configuration w and returns a real number.

For example, in a regression problem, the objective is to predict a continuous value, and a common loss function is the mean squared error (MSE), which measures the average of the squared differences between the actual values y_j associated with the record j and the corresponding predictions $\hat{y}_j(w)$, as follows:

$$L(w) = \frac{1}{N} \sum_{j=1}^{N} (y_j - \hat{y}_j(w))^2 \tag{1}$$

Instead, for a classification problem, the cross-entropy loss is often used. For a multiclass problem with M classes, it is expressed as:

$$L(w) = -\frac{1}{N} \sum_{j=1}^{N} \sum_{j'=1}^{M} y_{j,j'} \log(\hat{y}_{j,j'}(w)) \tag{2}$$

There are numerous other choices for $L(w)$, such as MAE [Mitchell, 1997], Huber Loss [Tong, 2023], KL divergence [Van Erven and Harremos, 2014], or Dice Loss [Li et al., 2019] (to name just a few). However, this document does not aim to provide a detailed discussion on the choice of the loss function; for a comprehensive review, refer to [Wang et al., 2020].

What is important to emphasize here is that training consists of adjusting the set of weights w in such a way as to minimize the error $L(w)$. In other words, the goal is to find the network weight configuration w^* such that:

$$w^* = \arg\min_w L(w) \tag{3}$$

ANNs excel in tasks such as image recognition [Quraishi et al., 2012], natural language understanding [Fanni et al., 2023], financial data analysis [Kurani et al., 2023], and even the creation of original content [Elbadawi et al., 2024], among many others. Their success lies in the fact that they are universal function approximators [Schäfer and Zimmermann, 2006]. Thus, whenever there is a process in which a certain input corresponds to a certain output, they aim to mimic and identify that process.

The fundamental concept behind neural networks is their ability to model complex relationships between data, making them a powerful tool for solving problems that are difficult to address with traditional algorithms.

Thanks to training techniques, they can learn autonomously from data, improving their performance as they receive more information [Haykin, 2009].

1.1 Problem presentation

What is typically done in equation (3) is adding a regularization term to the objective function that directs the network's training towards a weight configuration that satisfies not only the minimum of the function $L(w)$ but also some other property.

In this case, (3) becomes:

$$\min\{L(w) + \lambda\Omega(w) : w \in \mathbb{R}^n\} \tag{4}$$

Where λ is a hyperparameter, that is, a parameter that must be set before the model training and is not directly learned from the data. It is used to balance the terms $L(w)$ and $\Omega(w)$.

A typical example is ensuring that the weights do not overfit the instances seen by the network during training (training set) but rather generalize better in their predictions on unseen instances (test set). If we allowed the network weights to take any value, we might be able to correctly classify[2] the instances from the training set even in 100% of the cases, without actually learning the underlying structure of the data, and therefore without having learnt generalization capabilities.

For this reason, during the training of a neural network, it is undesirable for the weights to grow too much (in absolute value), and therefore a term is added in the problem (3) with the purpose of minimize the norm of the weights.

The most common choices to limit this growth of the weights are either:

$$\Omega(w) = ||w||_1 = \sum_{i=1}^{n} |w_i| \tag{5}$$

or:

$$\Omega(w) = ||w||_2 = \sum_{i=1}^{n} w_i^2 \tag{6}$$

We do not discuss the difference here but refer to [Melkumova and Shatskikh, 2017] for a detailed discussion.

Another example is structured pruning regularization [Cacciola et al., 2024], a regularization technique that promotes sparsity at the structural level (e.g., channels, filters, layers) in deep learning models, encouraging models that are more efficient in terms of memory and computation. There are several similar techniques that apply regularization to achieve analogous effects.

In this work, we choose to integrate a term into the network training problem that takes into

---

[2]This argument is not limited to classification problems.

account, in addition to the standard regularization to avoid overfitting, also the minimization of the entropy of the network's weights.

The weight tensor can indeed be seen as a sequence of values to which a metric, the entropy, can be associated, considering how well the space it occupies in memory lends itself to compression.

By training the network to minimize this metric, it is possible to obtain, once training is completed, a network that not only possesses good predictive capability for the problem at hand but is also more amenable to compression in terms of memory usage.

Achieving this result is crucial for integrating artificial intelligence algorithms into applications on devices with limited resources, such as smartphones, IoT devices, or embedded systems [Han et al., 2016].

Moreover, reducing the entropy of the weights can act as a form of regularization that limits the model's complexity, reducing the risk of overfitting. In other words, the network is "forced" to focus on the most relevant information rather than learning noise or non-general details.

1.1.1 Buckets Organization

To develop a training strategy while simultaneously compressing the neural network, we construct $\Omega(w)$ in such a way as to guide the training towards a quantization of the weights, i.e., towards a limited number of possible values for the weights.

Let C be the number of buckets into which we want to direct the weights, and let c_b be a variable that counts the number of weights that fall into bucket b, where $b = 0, \ldots, C-1$. Then, we can write the Shannon entropy as:

$$H = \sum_{b=0}^{C-1} c_b \log_2 c_b \tag{7}$$

The minus sign that is usually found in front of the entropy formula is actually conventional and serves to ensure that entropy is always a positive quantity. In fact, what is typically defined is the entropy over probabilities; these are values between 0 and 1, and thus their logarithm is always a negative quantity. Therefore, it is necessary to artificially insert a minus sign in front of the entropy summation to guarantee its positivity. However, in our case, c_b is a count, not a fraction, so we can define the entropy with the positive sign.


Figure 2: Buckets organization

The goal is to train the network so that the entropy is minimized. This will make the network more compressible once the weights are quantized.

Such quantization consists in assigning to all the weights within a specific bucket the central value of the bucket. If we assume that the weights are within a certain range $[\overline{w} - r, \overline{w} + r]^3$, the weights within the interval $[\overline{w} - r, \overline{w} - r + \frac{2r}{C})$ (i.e., the first bucket) will be assigned the central value $v_0 = \overline{w} - r + \frac{r}{C}$, those in the interval $[\overline{w} - r + \frac{2r}{C}, \overline{w} - r + 2\frac{2r}{C})$ (i.e., the second bucket) will be assigned the central value $v_1 = \overline{w} - r + \frac{3r}{C}$, and so on until the last bucket $[\overline{w} - r + (C-1)\frac{2r}{C}, \overline{w} + r]$ which will be assigned to the last value $v_{C-1} = \overline{w} + \frac{(C-1)r}{C}$.

The C buckets we have defined are equidistant from each other, with a width of $\frac{1}{C}$. In Figure 2, there is a pictorial representation of this arrangement.

Once the training aimed at minimizing entropy has been completed, and the trained weights have been quantized to the reference values, the sequence of weights can be treated as a sequence of symbols to be fed into compression algorithms, which can effectively encode it into a data structure that takes up less memory space.

1.1.2 METaQ Problem

In light of what has been presented in the previous sections, the problem of training the network with an entropy compression term (the so called *METaQ Problem*) can be written as follows:

$$\min\{L(w) + \lambda(\alpha R(w) + (1 - \alpha)\phi(w)) : \mathbb{R}^n\} \tag{8}$$

where λ is the hyperparameter introduced earlier that balances how much weight (in the sense of importance) to give to $L(w)$ and how much weight to give to the total regularization $\Omega(w)$ during training.

$\Omega(w)$ is composed of two parts, $R(w)$ and $\phi(w)$, whose relative importance is balanced by a second hyperparameter α, which expresses as a percentage how much $R(w)$ contributes to the total regularization $\Omega(w)$.

The function $R(w)$ represents the standard regularization term, while the function $\phi(w)$ represents the entropy regularization term, and is itself given by the following optimization problem (we indicate with $\mathfrak{B}$ the set of buckets):

---

[3]One could also think of $\overline{w} = \frac{1}{2}$ and $r = \frac{1}{2}$ to model the case where the weights are in $[0, 1]$, but for greater generality, we leave the interval in its general form. It is not necessarily the case that the most performant configuration of the network has weights in $[0, 1]$. In this special case we have that the vector of quantization is $v = [0, \frac{1}{C}, \frac{2}{C}, ..., \frac{C-1}{C}]$.

$$\phi(w) = \min \sum_{b \in \mathfrak{B}} c_b \log_2(c_b) \tag{9}$$

$$w_i = \sum_{b \in \mathfrak{B}} v_b x_{i,b} \qquad i \in I \tag{10}$$

$$\sum_{b \in \mathfrak{B}} x_{i,b} = 1 \qquad i \in I \tag{11}$$

$$c_b = \sum_{i \in I} x_{i,b} \qquad b \in \mathfrak{B} \tag{12}$$

$$x_{i,b} \in \{0,1\} \qquad i \in I, b \in \mathfrak{B} \tag{13}$$

As we previously said, in the problem (9)-(13), I is the set of weights in the network with cardinality $|I| = n$, $\mathfrak{B}$ is the set of buckets into which we want to direct the weights, v_b is the quantization value of bucket b, which, according to the construction made in Section 1.1.1, can be written as $v_b = \overline{w} - r + (2b+1)\frac{r}{C}$, where $b = 0, \ldots, C - 1^4$, $w_i \in [\overline{w} - r, \overline{w} + r]$ is the i-th weight of the network, and finally, the variables $x_{i,b}$ are binary variables defined as follows:

$$x_{i,b} = \begin{cases} 1, & \text{if weight } i \text{ belongs to bucket } b \\ 0, & \text{otherwise} \end{cases} \tag{14}$$

In optimization theory, the variables $x_{i,b}$ are called decision variables, and they play a crucial role since they are the true variables of optimization problems, including in the definition of $\phi(w)$.

From the definition of $x_{i,b}$, it follows that in the problem (9)-(13), the constraint (10), combined with the constraint (11), which imposes that only one of the $x_{i,b}$ is 1 in a given bucket, tells us that each weight value w_i in the network must equal exactly one of the representative values v_b of the buckets. The constraint (10) is important because it is the only point where the dependence of ϕ on w is explicitly stated.

The constraint (12), which defines c_b, simply imposes that c_b is the total number of weights associated with bucket b.

From the integrality of the $x_{i,b}$ expressed in (13), the integrality of c_b follows.

Thus, the combination of (10) and (11) defines the quantized representation of the weights, the constraint (12) prepares the data to compute the entropy, and the constraint (13) ensures the discrete nature of the quantization.

However, the problem (9)-(13) as defined involves a large number of binary decision variables $x_{i,b}$

---

[4]In the special case where $[\overline{w} - r, \overline{w} + r] \to [0,1]$, i.e., if $\overline{w} = \frac{1}{2}$, $r = \frac{1}{2}$, the first bucket goes from 0 to $\frac{1}{C}$ with a central value of $\frac{1}{2C}$, as also obtained from the expression $\overline{w} - r + (2b+1)\frac{r}{C}$ for $b = 0$.

and constraints. If we imagine an 8-bit quantization corresponding to $C = 256$ levels, we have up to $256 \times n$ decision variables, which, for a large-scale network, would translate into a number of constraints on the order of a billion, if not even more, making the problem intractable. Exactly solving a problem of this scale would require enormous computational power and would be impractical for real-world applications, probably not only for large neural networks but also for smaller networks.

Fortunately, this problem can be easily circumvented. In fact, to solve the training problem, what is necessary to provide to machine learning tools like `pytorch` is a gradient of the function we want to minimize. Typically, for standard loss functions and regularization terms, these gradients (i.e., $\nabla L(w)$ and $\nabla R(w)$) are well-known functions for the same tools. What we need to find is $\nabla \phi(w)$.

An important result from Lagrangian relaxation theory comes to our rescue: the optimal Lagrange multipliers β^* associated with the constraint (10) are precisely a (sub)gradient for $\phi(w)$ to feed into tools like `pytorch` to solve the METaQ problem.

The goal of this work is to illustrate how to find, through optimization techniques and algorithms, the vector β^* or at least an approximation $\overline{\beta}$ to solve the METaQ Problem in equation (8).

1.2 Literature review

We have seen the many advantages that neural network compression brings with it. Applications of these methods find ample space in so-called Large Language Models [Zhu et al., 2024], neural networks designed to understand and generate natural language, where the dimensionality of the models, as can be easily inferred from the name, is very large. The most famous ones, such as GPT4 [Achiam et al., 2023], LLaMA 3 [Dubey et al., 2024], PaLM [Chowdhery et al., 2023], and Claude [Enis and Hopkins, 2024], to name just a few, have hundreds of billions of parameters, which translates into a considerable dimensionality that is unlikely to have a direct application on client devices. This is why the largest companies in the sector are moving towards the implementation of compression strategies for these models [Knight, 2025], to ensure lower resource usage and energy consumption, not to mention that enabling execution on local devices increases user control over their data, improving privacy.

In the landscape of neural network compression algorithms, pruning algorithms certainly find ample space [Cheng et al., 2024]. Pruning removes non-essential weights or neurons in a neural network, thus reducing the overall number of parameters. It is based on the observation that many connections in over-sized models have a negligible impact on performance. The most common techniques are structured pruning (see for example [Santacroce et al., 2023] or [Ma et al., 2023]), which removes entire filters, channels, or layers, and unstructured pruning ([Frantar and Alistarh, 2023], [Syed et al., 2023], and [Zhang et al., 2023]), which removes individual weights leaving the network more irregular.

Another family of compression algorithms consists of the so-called quantization algorithms. Quan-

tization reduces the numerical precision of the network's weights and/or activations, for example, by switching from 32-bit floating-point numbers (FP32) to lighter representations such as INT8 or even binary. It is further divided into two main types: post-training quantization (PTQ), where quantization is applied after training (as in [Liu et al., 2021]), and quantization-aware training (QAT), which integrates quantization during training, allowing the network to adapt ([Liu et al., 2023], [Liu et al., 2022], [Ding et al., 2023], [Fu et al., 2023]). Naturally, all these types of algorithms can sometimes also be applied in pipelines, as in [Han et al., 2016].

From a quantization point of view, the strategy presented in this work can be placed somewhere between QAT and PTQ algorithms; indeed, during training, there is no true quantization. The introduction of the term (9) in the METaQ problem (the training problem of the network with the addition of the entropy regularization term) of equation (8) allows the network weights to *be directed* towards quantization, which will then be actually performed after training by replacing the weights with the reference value of the bucket assigned during training (and not by modifying the memory space occupied by the weights).

For the solution of the problem (9)-(13), [Pappalardo and Passacantando, 2012] was an indispensable guide for the foundations of optimization theory and operations research, while [Frangioni, 2005] provided the basis for the theory of Lagrangian relaxation.

The historical cornerstone [Kelley, 1960], together with the instructive application found in [Farruggia et al., 2019], allowed tackling the resolution of the Lagrangian relaxation subproblem of (9)-(13), providing the intuition for a specialized algorithm. Since it was possible to characterize the optimal solution of the relaxation in a simple way, it became feasible to efficiently determine the value of the Lagrangian solution as done in [Frangioni and Gorgone, 2013].

[D'Antonio and Frangioni, 2009], [Beck and Teboulle, 2009a], and [Frangioni, 2020] were of fundamental help in solving the dual problem; the first two with subgradient strategies and the latter as an invaluable guide in the sophisticated world of bundle methods.

Finally, for the network compression part, it is important to note that in this work we limited ourselves to considering, for simplicity's sake, zero-order entropy, meaning that each symbol of the source (the set of weights) was considered independent of the others, and no relationship or context between the symbols was taken into account. This is clearly a limitation, since in the neural network the weights are interconnected. For this reason, in addition to algorithms like Huffman coding and arithmetic coding, something like the FM-index [Ferragina et al., 2004] and its variants could have been used to exploit higher-order entropy k[5].

---

[5]Algorithms such as gzip-9 and zstd-22 have nonetheless been used, and in some tests, the latter has surpassed the 0-order entropy. See the results section for further details.

1.3   Thesis content

The organization of this document is as follows.

In Section 2, we will discuss the approximations made when handling the problem (9)-(13). We will then define a Lagrangian relaxation of the problem, which will allow us to decompose it into two solvable subproblems. We will discuss the analytical solution of the first and develop a custom algorithm to efficiently solve the second. Much of the discussion will focus on the latter point. At this stage, after stating and proving optimization theorems, we will use these results to solve the previously formulated Lagrangian relaxation through some appropriately discussed non-differentiable optimization algorithms. We will see that the solutions found in this way will also yield the sought approximation of β^* for $\nabla\phi$.

In Section 3, we will discuss how, once a low-entropy trained network is obtained, it can be compressed using well-known compression algorithms.

In Section 4, after introducing the well-known MNIST dataset and the famous LeNet-5 architecture specifically designed for its classification, we will present the results obtained by applying the strategy described and analyzed throughout this work.

Finally, in Section 5, we will discuss the implementation issues of METaQ, the precautions to take if one wishes to replicate the results on other networks, and potential future developments.

Section 6 contains acknowledgements, while Appendix 1 will summarize all the hyperparameters present in METaQ with a brief descriptions and their configuration used to achieve the results for the best model.

2   LAGRANGIAN RELAXATION

The first thing to do in order to handle (9)-(13) is to ensure that $\phi(w)$ is at least well-defined everywhere. For now, in fact, $\phi(w)$ is not even continuous, since the right-hand side of (10) spans a discrete set of values, while the left-hand side spans a continuous one. The problem is that $x_{i,b}$ is a binary variable, which makes the problem discontinuous. Therefore, the first thing to do is to relax (13) to:

$$x_{i,b} \in [0,1] \qquad i \in I, b \in \mathfrak{B} \tag{15}$$

In the problem (9)-(13), the constraints (12) tie together the x corresponding to different weights. If it were possible to decouple the elements of c from the elements of x, the source of non-linearity (i.e., the entropy) could be eliminated.

This is possible through Lagrangian relaxation techniques. What is done, in fact, is to introduce the Lagrange multipliers $\{\xi_b\}_{b\in\mathfrak{B}}$, each of which is associated with one of the constraints in (12). The goal at this point is to solve a problem in which these constraints do not need to be satisfied exactly (hence the term "relaxed"), but rather they should be satisfied "as much as possible." Once the LHS and RHS are moved to the same side, the constraints are moved into the objective function, and we minimize, along with what was previously to be minimized, the deviation of this difference from 0. The solution that will be found will therefore be one that drives this term toward 0, i.e., towards satisfying the relaxed constraints. The problem thus becomes one of finding the function (concave in the variables ξ):

$$\phi_w(\xi) = \min\left\{\sum_{b\in\mathfrak{B}} c_b \log_2(c_b) + \sum_{b\in\mathfrak{B}} \xi_b(\sum_{i\in I} x_{i,b} - c_b) : (10), (11), (15)\right\} \tag{16}$$

in order to maximize it:

$$\max\{\phi_w(\xi) : \xi \in \mathbb{R}^C\} \tag{17}$$

In correspondence with the solution ξ^* of (17), the corresponding optimal multipliers β^* associated with the constraints (10) (or more precisely, an approximation $\overline{\beta}$ of them) will have been found, and thus a (sub)gradient $\nabla\phi(w)$ for $\phi(w)$.

In fact, holds the following:

Theorem 1. *Let $\phi(w) = \min\{f(x) : g(x,w) = 0, x \in X\}$, where $g(x,w) = w - \sum_{b\in\mathfrak{B}} v_b x_{i,b}$ as in Equation (10), and X constraints x to stay in a certain region of space (it represents the other contraints stated in (11) and (12)).*

Then, the optimal Lagrange multipliers β^* associated with the constraints $g(x,w) = 0$ are a sub-gradient for the function $\phi(w)$.

Proof. The Lagrangian function is:

$$L(x,w,\beta) = f(x) + \beta^t \cdot g(x,w) \tag{18}$$

Therefore we can write $\phi(w)$ as:

$$\phi(w) = \min\{f(x) : g(x,w) = 0, x \in X\} = \min\{L(x,w,\beta)\} \tag{19}$$

Let $x^*(w)$ be the optimal value for ϕ in w, that is:

$$\phi(w) = L(x^*(w), w, \beta^*(w)) = f(x^*(w)) + \beta^{*t} \cdot g(x^*(w), w) \tag{20}$$

Since $x^*(w)$ is optimal, we have $g(x^*(w), w) = 0$, hence $\phi(w) = f(x^*(w))$. Now, $\forall w'$ we have[6]:

$$g(x, w') = w' - \sum_{b \in \mathfrak{B}} v_b x_{i,b} = g(x, w) + w' - w \tag{21}$$

and evaluating this expression in $x^*(w)$:

$$g(x^*(w), w') = g(x^*(w), w) + w' - w = w' - w \tag{22}$$

For $\phi(w')$ we certainly have:

$$\phi(w') \geq f(x^*(w)) + \beta^{*t} \cdot g(x^*(w), w') \tag{23}$$

Therefore, by substituting $f(x^*(w))$ with $\phi(w)$, and $g(x^*(w), w')$ with $w' - w$, we finally have:

$$\phi(w') \geq \phi(w) + \beta^{*t} \cdot (w' - w) \tag{24}$$

That is β^* is a sub-gradient for $\phi(w)$.

$\square$

2.1 Analitic subproblem resolution

In this section, we solve the subproblem of (16) related to the variables c_b, that is:

$$\min\{c_b \log_2 c_b - \xi_b c_b\} \tag{25}$$

Proposition 1. *The analytical solution of* (25) *is* $c_b^* = e^{2\xi_b - 1}$.

Proof. The optimality condition is written as:

$$\frac{d}{dc_b}(c_b \log_2(c_b) - \xi_b c_b) = 0 \implies \frac{d}{dc_b}\left(c_b \frac{\log(c_b)}{\log(2)} - \xi_b c_b\right) = 0 \tag{26}$$

Computing the derivative w.r.t. c_b we obtain:

---

[6]Here, we rely on the fundamental assumption regarding the form of the constraints. It is important to note that if the assumption of constraint linearity were to fail, the shape of the subgradient would become significantly more complex.

$$\frac{1}{\log(2)} \left(\log(c_b^*) + 1\right) - \xi_b = 0 \implies \log(c_b^*) = \log(2)\xi_b - 1 \tag{27}$$

And finally:

$$\log_2(c_b^*) = \frac{\log(2)\xi_b - 1}{\log(2)} = \xi_b - \frac{1}{\log(2)} \tag{28}$$

which results in $c_b^* = 2^{\xi_b - \frac{1}{\log(2)}}$ or equivalently $c_b^* = e^{2\xi_b - 1}$.

$\square$

As can be observed, the dependence of c_b^* on the multipliers ξ_b is exponential, which is why it is more appropriate to transform the problem (25) into:

$$\min\{c_b \log_2 c_b - \xi_b c_b : c_b \in [\underline{c_b}, \overline{c_b}]\} \tag{29}$$

That is, truncate the solutions to a certain interval; those for which $c_b > \overline{c_b}$ are set to $\overline{c_b}$, while those for which $c_b < \underline{c_b}$ are set to $\underline{c_b}$.

The most natural choice for $\overline{c_b}$ is $\overline{c_b} = n$, since no more than n weights can end up in each bucket (this configuration represents the situation in which all the weights are in bucket b).

The choice for $\underline{c_b}$ is somewhat more delicate. Since c_b must definitely be a positive value as it represents the count of weights that fall into bucket b, we need to impose at least $\underline{c_b} \geq 0$.

Furthermore, it is certainly necessary to ensure that the logarithm exists, so the condition is stronger and $\underline{c_b} > 0$ must hold.

However, one must also avoid setting this value too close to 0, as it could significantly influence the optimal value of the objective function.

It should be a value that is not too close to 0, so as not to overly influence the overall entropy, and not too large to avoid truncating all solutions to it.

In our case, as we will explore in more detail in the results section, with a choice of $\underline{c_b} = 0.01$, the entropy of the network we analyzed is not significantly affected.

For this reason, this value could serve as a starting point for more precise tuning. It would definitely be more appropriate to leave it as a hyperparameter to be chosen based on the type of network and the results obtained after various trials.

2.2 Knapsack subproblem

In this section, we discuss the solution of the subproblem of (16) related to the decision variables x. The problem is as follows:

$$\min_x \sum_{b \in \mathfrak{B}} \xi_b x_{i,b} \tag{30}$$

$$w_i = \sum_{b \in \mathfrak{B}} v_b x_{i,b} \tag{31}$$

$$\sum_{b \in \mathfrak{B}} x_{i,b} = 1 \tag{32}$$

$$x_{i,b} = [0,1] \tag{33}$$

And it has to be solved for each weight $w_i, i \in I$.

In the following, we will omit indexing on i since these n problems refer to weights that are independent of each other, and there is therefore no risk of ambiguity. We will write the sums in terms of dot products wherever possible.

The given problem has strong similarities to the knapsack problem, especially to the fractional variant. The goal is indeed to minimize or maximize a linear function involving the cost or value of the items, with constraints on the combinations of selected items.

The variables $x_{i,b}$ can be interpreted as a measure of how much to "include" each item. The main difference is the unique selection constraint (32), which does not appear in the classical version of the knapsack problem.

For solving the problem (30)-(33), we will assume that the vector ξ and the vector v are ordered. For the latter, the ordering property follows immediately from its construction, while for the former, it will need to be ensured at the implementation level. As we shall see, this sorting operation only needs to be performed once for all weights during the training phase.

2.2.1 CVXPY Solution

One way to solve the problem (30)-(33) is to rely on an optimization solver. In this work, we have chosen CVXPY, a Python library for modeling and solving mathematical optimization problems, particularly convex optimization problems.

CVXPY does not directly implement solving algorithms but relies on external solvers to solve optimization problems, which in turn use algorithms such as the simplex method (and variants), interior-point methods (and variants), etc. We refer to the official documentation for details [cvxpy.org].

The key point here is that this library solves the problem (30)-(33) but does not do so in the most efficient possible way. Indeed, since it is designed to solve a wide range of optimization problems, it does not provide specialized tools for a specific problem (as it clearly should).

For this reason, it was considered necessary, for the success of this work, to implement a specific algorithm for our needs, that is, to solve the problem (30)-(33).

Once this algorithm is presented, we will compare execution times for solving (30)-(33) to assess the effectiveness of implementing a specialized method.

2.2.2 Cutting Plane Algorithm

The starting point for implementing a specialized algorithm to solve (30)-(33) is Kelley's Cutting Plane algorithm [Kelley, 1960].

This algorithm is an iterative method used to solve convex optimization problems, specifically for approximating a convex programming problem through a finite number of iterations. It is one of the "cutting plane" methods, which aim to progressively narrow down the search domain to converge toward an optimal solution.

The fundamental idea is to approximate the objective function using linear underestimations (tangent supports), relying on the convexity properties of the function.

In our case, let (P) be the following primal problem:[7]

$$\min_{x} \left\{ \xi^t \cdot x : w = v \cdot x, \sum_{b=0}^{C-1} x_b = 1, x_b \in [0,1], \ \forall b = 0, \dots, C-1 \right\} \tag{P}$$

Relaxing the constraint $w = v^t \cdot x$, yields a Lagrangian:

$$L(x, \lambda) = \xi^t \cdot x + \lambda(w - v^t \cdot x) \tag{34}$$

And therefore, we can define a dual function:

$$\phi(\lambda) = \min_{x \in X} L(x, \lambda) \tag{35}$$

where $X = \{x \in [0,1]^C : \sum_{b=0}^{C-1} x_b = 1\}$.

---

[7]The formulation is entirely equivalent to the problem (30)-(33) where the summation has been replaced with the dot product.

Proposition 2. *If x_1, $x_2 \in X$ then $\forall \theta \in [0,1]$, $\theta x_1 + (1-\theta)x_2 \in X$*

Proof. What we have is that:

$$\theta x_1 + (1-\theta)x_2 \in [0,1]^C \tag{36}$$

and this is because for each component:

$$(\theta x_1 + (1-\theta)x_2)_b = \theta x_{1b} + (1-\theta)x_{2b} \in [0,1] \tag{37}$$

Indeed, since $\theta x_{1b} + (1-\theta)x_{2b}$ is the sum of two non-negative terms, it will certainly be non-negative.

Furthermore:

$$\theta x_{1b} \leq \theta \cdot 1 = \theta \tag{38}$$

and

$$(1-\theta)x_{2b} \leq (1-\theta) \cdot 1 = (1-\theta) \tag{39}$$

Therefore, summing the two inequalities, we obtain:

$$\theta x_{1b} + (1-\theta)x_{2b} \leq \theta + (1-\theta) = 1 \tag{40}$$

i.e.

$$\theta x_{1b} + (1-\theta)x_{2b} \in [0,1] \tag{41}$$

On the other hand:

$$\sum_{b=0}^{C-1} \theta x_{1b} + (1-\theta)x_{2b} = \theta \sum_{b=0}^{C-1} x_{1b} + (1-\theta) \sum_{b=0}^{C-1} x_{2b} = \theta + (1-\theta) = 1 \tag{42}$$

And therefore:

$$\theta x_1 + (1-\theta)x_2 \in X \tag{43}$$

$\square$

The Dual Problem (D) is:

16

$$\max_{\lambda \in \mathbb{R}} \phi(\lambda) \tag{D}$$

To find a solution of (D) we use the cutting plane in the following way.

1) We initialize $x', x'' \in X$ such that $w \geq v^t \cdot x'$ and $w < v^t \cdot x''$. (For example, we can initialize x' as the vector with C components, all 0 except for a 1 in the component b' corresponding to the largest $v_{b'} \leq w$, and x'' as the vector with C components, all 0 except for a 1 in the component b'' corresponding to the smallest $v_{b''} > w$). At the implementation level, the assumption of ordering for v is crucial for identifying b' and b''.

2) At this point, since (D) is equivalent to (see [Farruggia et al., 2019], [Frangioni, 2005]):

$$\max_{\phi, \lambda} \{\phi : \phi \leq L(x, \lambda), x \in X\} \tag{44}$$

and since there are only x' and x'' in X, there are only 2 constraints, namely:

$$\phi \leq L(x', \lambda) \text{ and } \phi \leq L(x'', \lambda) \tag{45}$$

Let (λ^+, ϕ^+) the intersection between the two curves:

$$\phi = L(x', \lambda) = \xi^t \cdot x' + \lambda(w - v^t \cdot x') \tag{46}$$

$$\phi = L(x'', \lambda) = \xi^t \cdot x'' + \lambda(w - v^t \cdot x'') \tag{47}$$

in the (λ, ϕ) plane. We are certain that such an intersection exists because one curve has a positive slope and the other has a negative slope. By subtracting the second equation from the first and isolating λ^+, what we obtain is:

$$\lambda^+ = \frac{\xi^t \cdot (x' - x'')}{(v^t \cdot x' - w)(v^t \cdot x'' - w)} \tag{48}$$

and

$$\phi^+ = L(x', \lambda^+) = \xi^t \cdot x' + \lambda^+(w - v^t \cdot x') \tag{49}$$

or equivalently $(L(x', \lambda) = L(x'', \lambda))$:

$$\phi^+ = L(x'', \lambda^+) = \xi^t \cdot x'' + \lambda^+(w - v^t \cdot x'') \tag{50}$$

Let x^+ be the solution of the relaxed problem, i.e.,

$$x^+ = \operatorname*{argmin}_{x \in X}(\xi - \lambda^+ v)^t \cdot x \tag{51}$$

and since ξ (just like v) is ordered in increasing order, x^+ is a vector with 0 everywhere except for a 1 in the component where $\xi_b - \lambda v_b$ is minimum (i.e., for $b = 1$).

Case $\phi(\lambda^+) = \phi^+$.

When $\phi(\lambda^+) = \phi^+$, the point (λ^+, ϕ^+) lies exactly on the dual function $\phi(\lambda)$, i.e.,

$$\phi^+ = \min_{x \in X}(\xi^t \cdot x + \lambda^+(w - v^t \cdot x)) \tag{52}$$

By the Strong Duality Theorem, the optimal value of the primal problem (P) coincides with the optimal value of the dual problem (D):

$$\min_{x \in [0,1]^C} \xi^t \cdot x = \max_{\lambda \in \mathbb{R}} \phi(\lambda) = \phi^+ \tag{53}$$

Let us now define $x(\theta)$ as any convex combination of x' and x'':

$$x(\theta) = \theta x' + (1 - \theta)x'' \tag{54}$$

e let θ^* such that

$$v^t \cdot x(\theta^*) = w \tag{55}$$

Proposition 3. *$x(\theta^*)$ is a solution for (P).*

Proof. We need to prove that $x(\theta^*)$ satisfies all the primal constraints and that the optimal value of the objective function (i.e., ϕ^+ in this case) coincides with $\xi^t \cdot x(\theta^*)$.

As for the constraints, $x(\theta^*)$ satisfies (31) by construction, and satisfies (32) and (33) by Proposition 2.

As for the value of the objective function at $x(\theta^*)$, we have:

$$\xi^t \cdot x(\theta^*) = \xi^t(\theta^* x' + (1 - \theta^*)x'') = \theta^* \xi^t \cdot x' + (1 - \theta^*)\xi^t \cdot x'' \tag{56}$$

Now, since:

$$\phi^+ = \xi^t \cdot x' + \lambda^+(w - v^t \cdot x') \tag{57}$$

or equivalently:

$$\phi^+ = \xi^t \cdot x'' + \lambda^+(w - v^t \cdot x'') \tag{58}$$

we can write:

$$\phi^+ = \xi^t \cdot x^i + \lambda^+(w - v^t \cdot x^i), \ \text{with} \ i =','' \tag{59}$$

from which we can obtain:

$$\xi^t \cdot x^i = \phi^+ - \lambda^+(w - v^t \cdot x^i), \ \text{with} \ i =','' \tag{60}$$

hence, by substituting in (56):

$$\xi^t \cdot x(\theta^*) = \theta^*(\phi^+ - \lambda^+(w - v^t \cdot x')) + (1 - \theta^*)(\phi^+ - \lambda^+(w - v^t \cdot x'')) \tag{61}$$

By doing some calculations, we can arrive at the expression:

$$\xi^t \cdot x(\theta^*) = \phi^+ - \lambda^+[\theta^*(w - v^t \cdot x') + (1 - \theta^*)(w - v^t \cdot x'')] \tag{62}$$

But:

$$0 = w - v^t \cdot x(\theta^*) = w - (\theta^* v^t \cdot x' + (1 - \theta^*)v^t \cdot x'') = \theta^*(w - v^t \cdot x') + (1 - \theta^*)(w - v^t \cdot x'') \tag{63}$$

Therefore:

$$\xi^t \cdot x(\theta^*) = \phi^+ \tag{64}$$

$\square$

19

Case $\phi(\lambda^+) \neq \phi^+$ and $w = v^t \cdot x^+$.

In this case, x^+ minimizes $L(x, \lambda^+)$ (see Equation (23)), but since it satisfies constraint (2) (as well as constraints (3) and (4)), it is also optimal with optimal value:

$$\phi(\lambda^+) = \xi^t \cdot x^+ + \lambda^+(w - v^t \cdot x^+) = \xi^t \cdot x^+ \tag{65}$$

Case $w \neq v^t \cdot x^+$.

If $w > v^t \cdot x^+$, then x^+ does not yet satisfy constraint (2). Since the value of $v^t \cdot x^+$ is too small compared to w, it means that we need to create a new point x that gives a value of $v^t \cdot x$ closer to w. Consequently, we update the set X for the next iteration of the cutting plane algorithm by replacing x' with x^+ because x^+ is closer to the constraint than x', yet still maintains $v^t \cdot x^+ < w$. The new set X becomes $X = \{x^+, x''\}$, and the procedure is repeated.

If $w < v^t \cdot x^+$, then $v^t \cdot x^+$ is too large compared to w, so we replace x'' with x^+ since x^+ is closer to the constraint than x'', but still maintains $v^t \cdot x^+ > w$. The new set X becomes $X = \{x', x^+\}$.

2.2.3 Specialized Algorithm

Finally, we are ready to construct a specific algorithm for solving the problem (30)-(33).

The basic idea is that when $\lambda \to -\infty$, the solution[8]:

$$x^* = \arg\min_{x \in X}(\xi - \lambda v)^t \cdot x \tag{66}$$

The active component[9] of the solution is only the first component. This is because, as we stated earlier, both ξ and v are ordered in increasing order, and thus the minimum of the vector $\xi - \lambda v$ corresponds to the minimum of the two vectors, that is, the first component.

Conversely, if $\lambda \to +\infty$, the solution to (66) has as the active component only the last component.

As λ moves from $-\infty$ to $+\infty$, the minimum of the vector $\xi - \lambda v$ moves to the right, transitioning from $b = 0$ to $b = C - 1$.

We will call *breakpoints* the values of b where the minimum of $\xi - \lambda v$ switches from one component to another as λ moves from $-\infty$ to $+\infty$, and we will denote with $\overline{B}$ that C-dimensional vector with zeros everywhere and ones at the breakpoints. These breakpoints occur at the values λ^* such that $\xi_{b'} - \lambda^* v_{b'} = \xi_{b''} - \lambda^* v_{b''}$, where $b', b'' = 0, ..., N_B - 1$, and N_B is the number of breakpoints. It is

---

[8]See equation (51) where λ is left as a free variable.
[9]By active components of a vector, we mean its non-zero components.

evident that $N_B \leq C$ since at most we can have C breakpoints. Solving for λ^* gives:

$$\lambda^* = \frac{\xi_{b'} - \xi_{b''}}{v_{b'} - v_{b''}} \tag{67}$$

We have mentioned that two breakpoints are certainly $b = 0$ and $b = C - 1$; now we need to determine what the others are in order to construct $\overline{B}$. To find them, we start with the first one ($b = 0$) and calculate the minimum of $\frac{\xi_b - \xi_0}{v_b - v_0}$ for $b = 1, \ldots, C - 1$. In other words, the minimum is calculated among the elements to the right of the last breakpoint found.

In Figure 3, we present a depiction of what a possible configuration of the breakpoints vector might look like.

$$[1,0,0,1,1,0,1,0,0,0,1,1,0,1,1] \quad (C = 15)$$

Breakpoints

Figure 3: An exemplary depiction of a possible configuration for the breakpoints vector with $C = 15$.

We observe that the vector $\overline{B}$ depends exclusively on ξ and not on the value of w (not even on the vector v that is fixed throughout the training). This property can be exploited to implement the entire algorithm in a parallelized manner instead of solving it for each w_i sequentially at each training step (which would result in a long time of execution).

Once the vector $\overline{B}$ is constructed, the solution to the problem (30)-(33) is determined by w_i.

Indeed, depending on the value of w_i, we can compute the vectors that in the cutting-plane algorithm were represented by x' and x''. These vectors each have only one active component: in the first case, corresponding to the largest breakpoint b' such that $v_{b'} \leq w_i$, and in the second case, to the largest breakpoint b'' such that $v_{b''} \geq w_i$. Once b' and b'' are found, the convex combination of x' and x'' is taken, as in the cutting-plane algorithm (see Equation (55)). Optimal lagrangian multipliers are given by equation (67).

The edge cases where $w_i \leq v_0$ and $w_i \geq v_{C-1}$ correspond to solutions with a single active component, at $b = 0$ and $b = C - 1$, respectively.

2.2.4   Algorithms comparison

The first test conducted was to compare the speed of the cutting plane algorithm against that of the CVXPY library. Using $C = 256$, 1000 instances of the problem (30)-(33) were run, and the execution time for the cutting plane algorithm was $0.20 sec$, while for the CVXPY library it was $3.19 sec$, resulting in an efficiency improvement of a factor of around $16\times$.

For a comparison with the vectorized version of the cutting plane algorithm, 10000 instances of the problem (30)-(33) were run with $C = 5$, achieving an efficiency gain in terms of execution time of more than a factor of $2378\times$.

This result demonstrates how necessary it was to implement an alternative strategy to the standard CVXPY approach for solving the problem (30)-(33).

Finally, a comparison was also made between the vectorized cutting plane algorithm and the specialized algorithm (also in its vectorized version). Since both are fast algorithms, 44000 instances of the problem were run 1000 times (this value was chosen because it is of the order of the weights in LeNet-5 - see section of results) with $C = 256$. The execution time for the cutting plane algorithm was $215.98 sec$, while the specialized algorithm took $18.66 sec$, resulting in an efficiency gain of a factor of $11.58\times$.

The merit of these results can undoubtedly be attributed primarily to the structure of the refined algorithms discussed throughout this work; however, a crucial role was played by the ability to apply parallelization of the computations, enabling the simultaneous execution of multiple instances of the problem (30)-(33). In fact, for the specialized algorithm, the bulk of the complexity lies in constructing the vector $\overline{B}$; however, $\overline{B}$ is common to all w_i, so, provided we can run multiple instances of the problem for the different w_i simultaneously (which PyTorch handles very well as it was specifically designed for solving tensor-based problems), most of the work is done only once.

Next, we will discuss the complexity of the specialized algorithm.

As mentioned, the most computationally expensive part is calculating $\overline{B}$. Fortunately, this only needs to be done once for all the weights; however, this calculation involves finding a minimum that can potentially happen C times. The minimum is taken over a vector that gets progressively smaller, but still, the number of operations performed is $O(C^2)$. The worst-case scenario occurs when the breakpoints are exactly C, and thus for each breakpoint, a minimum of length $O(C)$ must be calculated.

We now provide a heuristic argument showing that the worst-case scenario practically never occurs if C is "large enough," also giving a characterization of the term "large enough".

To do this, it is sufficient to calculate the (empirical) probability of having a certain number of breakpoints.

In Figure 4, we can see that for $C = 256$, meaning a $\overline{B}$ with 256 entries, out of 10,000 instances of the problem, no more than 13 breakpoints are ever obtained. Moreover, the probability seems to be modeled by a Beta distribution, peaking around a very low value for the most probable number of breakpoints.

By fitting the empirical data to a beta distribution the results is what is shown Figure 5.

With this analytical distribution, it is possible to estimate the probability of obtaining a certain

number of breakpoints. In particular, for this Beta distribution, the probability of having 256 breakpoints in a vector $\overline{B}$ of 256 elements is $1.82 \cdot 10^{-157}$.


Figure 4: Histogram of the occurrences of having a certain number of breakpoints for a $\overline{B}$ with $C = 256$.
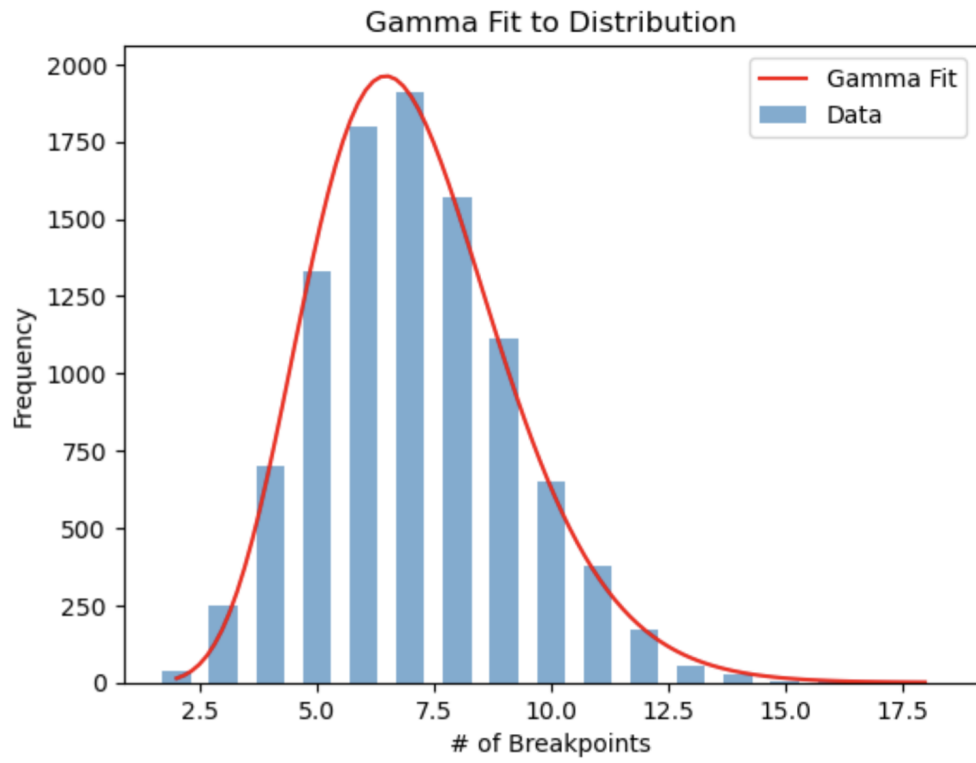

Figure 5: Fit of a Beta distribution to the empirical data.

2.3 Lagrangian Dual solution

At this point, at a fixed step of the training cycle we have found, for each weight of the network, the optimal value for c_b (which we denoted as c_b^*) and the optimal value of the decision variables $x_{i,b}$ (which we denoted as $x_{i,b}^*$). Corresponding to these values, we have therefore fully defined the function $\phi_w(\xi)$, that is:

$$\phi_w(\xi) = \sum_b c_b^* \log_2(c_b^*) + \sum_b \xi_b \left(\sum_i x_{i,b}^* - c_b^* \right) \tag{68}$$

This function must be maximized in order to solve the Lagrangian dual (17), which we present here for the sake of readability.

$$\max\{\phi_w(\xi) : \xi \in \mathbb{R}^C\} \tag{69}$$

Below, we present some methods through which this problem can be solved. We focus on two families: the so-called subgradient methods and the so-called bundle methods.

With reference to equation (68), we state the following theorem:

Theorem 2. $g = \{\sum_{i=1}^n x_{i,b}^* - c_b^*\}_{b=0}^{C-1}$ *is a super-gradient for* $\phi_w(\xi)$

Proof. From Equation (68), it follows that:

$$\sum_b c_b^* \log_2(c_b^*) = \phi_w(\xi) - \sum_b \xi_b \left(\sum_i x_{i,b}^* - c_b^* \right) \tag{70}$$

For any ξ', we have:

$$\phi_w(\xi') = \min\left(\sum_b c_b \log_2(c_b) + \sum_b \xi_b' \left(\sum_i x_{i,b} - c_b \right) \right) \tag{71}$$

And evaluating this expression at the optimum for ξ, we get:

$$\phi_w(\xi') \leq \sum_b c_b^* \log_2(c_b^*) + \sum_b \xi_b' \left(\sum_i x_{i,b}^* - c_b^* \right) = \phi_w(\xi) - \sum_b \xi_b \left(\sum_i x_{i,b}^* - c_b^* \right) + \sum_b \xi_b' \left(\sum_i x_{i,b}^* - c_b^* \right) \tag{72}$$

that is

$$\phi_w(\xi') \leq \phi_w(\xi) + \sum_b (\xi_b' - \xi_b) \sum_i (x_{i,b}^* - c_b^*), \ \forall \xi' \tag{73}$$

i.e. $\{\sum_i x_{i,b}^* - c_b^*\}_{b=0}^{C-1}$ is a supergradient for $\phi_w(\xi)$.

$\square$

2.3.1 Supergradient Methods

To solve (17), we follow a strategy similar to the one used for differentiable problems, that is, we construct an iterative algorithm where at each iteration k the variable ξ (which we can index as ξ_k) is updated by moving it along the direction of the gradient (supergradient in our case) of the function. Let γ_k denote the step size by which this movement occurs along the gradient direction. We can then write the update for the next step as follows:

$$\xi_{k+1} = \xi_k + \gamma_k g_k \tag{74}$$

The choice of γ_k is crucial, and the simplest way to choose it is using the so-called diminishing square-summable (DSS) stepsize, that is,

$$\gamma_k = \frac{1}{k} \tag{75}$$

Proposition 4. *With the choice* (75)*, the iterative algorithm* (74) *converges.*

Proof. We have:

$$||\xi_{k+1} - \xi^*||^2 = ||\xi_k + \gamma_k g_k - \xi^*||^2 = ||\xi_k - \xi^*||^2 + (\gamma_k)^2||g_k||^2 + 2\langle \xi_k - \xi^*, \gamma_k g_k \rangle \tag{76}$$

which, by factoring out γ_k from the inner product, can be written as:

$$||\xi_{k+1} - \xi^*||^2 = ||\xi_k - \xi^*||^2 + (\gamma_k)^2||g_k||^2 + 2\gamma_k \langle \xi_k - \xi^*, g_k \rangle \tag{77}$$

Now, by the definition of a supergradient, we have:

$$f(\xi^*) \leq f(\xi_k) + \langle \xi^* - \xi_k, g_k \rangle \Rightarrow f(\xi^*) - f(\xi_k) \leq \langle \xi^* - \xi_k, g_k \rangle \tag{78}$$

That is, by changing the sign to make it explicit, we get $\langle \xi_k - \xi^*, g_k \rangle$:

$$\langle \xi^* - \xi_k, g_k \rangle \geq f(\xi^*) - f(\xi_k) \Rightarrow \langle \xi_k - \xi^*, g_k \rangle \leq f(\xi_k) - f(\xi^*) \leq 0 \tag{79}$$

In the last inequality, we used the fact that the optimal value $f(\xi^*)$ is a maximum. Therefore, what we have is that:

$$||\xi_{k+1} - \xi^*||^2 \leq ||\xi_k - \xi^*||^2 + (\gamma_k)^2 ||g_k||^2 + 2\gamma_k(f(\xi_k) - f(\xi^*)) \tag{80}$$

Therefore, by choosing γ_k as in equation (75), the term $(\gamma_k)^2 ||g_k||^2$ becomes negligible for large k compared to the term $2\gamma_k(f(\xi_k) - f(\xi^*))$, and since the latter term is negative, we have that with this choice of γ_k,

$$||\xi_{k+1} - \xi^*||^2 \leq ||\xi_k - \xi^*||^2 \tag{81}$$

That is, the iterates ξ_k converge to the optimal ξ^*.

$\square$

Given its simplicity in implementation, the DSS stepsize method is very common in optimization practices. However, it has some limitations; it converges slowly and does not take into account information from previous iterations.

A more often considered alternative, which is generally more efficient, is the FISTA (Fast Iterative Shrinkage-Thresholding Algorithm) method, first introduced in 2009 by Amir Beck and Marc Teboulle in [Beck and Teboulle, 2009b]. FISTA is an evolution of gradient-based methods designed to accelerate convergence compared to the basic gradient method or methods with DSS stepsize. It combines the gradient method with an acceleration term based on an inertial term inspired by Nesterov's techniques [Nesterov, 1983]. This term takes into account the "history" of previous iterations, helping to push the solution more quickly towards the optimal minimum.

Introducing a hyperparameter ζ_{FISTA} as the convergence step, an auxiliary variable is used.

$$\sigma_k = \xi_k + \frac{t_{k-1} - 1}{t_k}(\xi_k - \xi_{k-1}) \tag{82}$$

where t_k is an increasing sequence that accelerates convergence, such as:

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2} \tag{83}$$

and the iterate ξ_k is updated using the auxiliary variable σ_k as follows:

$$\xi_{k+1} = \sigma_k + \frac{1}{\zeta_{FISTA}}g_k \tag{84}$$

2.3.2 Bundle Methods

In this section, we discuss bundle methods, a class of optimization algorithms (both linear and nonlinear) widely used to tackle non-differentiable problems or those with objective functions that exhibit non-smooth behaviors. These methods are a generalization of subgradient methods, improving their stability and convergence speed.

The main idea behind bundle methods is to construct an approximation of the objective function based on a set $B^{(k)}$ (called the bundle) of local information, such as function evaluations and subgradients, which are then used to iteratively generate a direction of improvement. This approach proves particularly useful in convex (or concave) non-differentiable optimization problems, where the objective function is not sufficiently smooth for classical gradient-based methods to be applied.

In this family of methods, we will discuss the Proximal Bundle Method (PBM), a variant of bundle methods that introduces a regularization term to make the optimization more stable and effectively handle non-differentiability. The main idea is to minimize an aggregated approximation of the objective function, balanced with a penalization term that accounts for the distance between the current point and the next candidate point.

The PBM approach is based on the iterative solution of the regularized problem:

$$\xi^{(k+1)} = \arg\max_{\xi}(\phi_w^{(k)}(\xi) - \frac{\zeta_{PBM}}{2}||\xi - \xi^{(k)}||^2) \tag{85}$$

where

- $\underline{\phi}_w^{(k)}(\xi)$ is a modeled underestimation of the objective function, constructed using the bundle (supergradients and function values at previous points);

- $||\xi - \xi^{(k)}||^2$ is a penalty term that encourages iterations to stay close to the current iterate $\xi^{(k)}$;

- $\zeta_{PBM} > 0$ is a regularization parameter that balances the weight between the function approximation and the penalty term.

At each iteration, the supergradients $g^{(k)}$ and the function values $\phi_w(\xi^{(k)})$ are collected at a series of $B_s = |B^{(k)}|$ previous points $\xi_j^{(k)}$, $j = 1, \ldots, B_s$.

The parameter B_s represents the bundle size and is a key hyperparameter of the algorithm. If the bundle size is too small, the algorithm may not have enough historical information to construct an accurate approximation of the objective function, risking suboptimal choices. In such cases, bundle methods behave similarly to subgradient methods, losing the advantage of constructing the bundle. On the other hand, if the bundle size is too large, the algorithm may become more complex and computationally expensive because the number of points to be considered grows.

Thus, at each iteration k, the goal is to solve the subproblem among the elements of the bundle $\{\xi_j^{(k)}, \phi_{w,j}^{(k)}, g_j^{(k)}\}_{j=1}^{B_s}$

$$j^* = \arg\max_{j \in B^{(k)}} \{\phi_{w,j}^{(k)} + g_j^{(k)t} \cdot (\xi - \xi_j^{(k)})\} \tag{86}$$

and the update of the iterates is performed through:

$$\xi^{(k+1)} = \xi_{j^*}^{(k)} + \frac{1}{\zeta_{PBM}} g_{j^*}^{(k)} \tag{87}$$

3 NETWORK COMPRESSION

Once the problem (17) has been solved, or at least an approximate solution $\overline{\xi}$ has been found, we have also solved, in correspondence with this solution, the problem (30)-(33), and thus we have found the corresponding optimal multipliers β^* (or at least their approximation $\overline{\beta}$) of constraints (31), which, as we have seen, correspond to the subgradient of $\phi(w)$ to be used in training the network and thus solving the METaQ Problem (8).

We remind that this training will aim to push the network weights w not only towards minimizing the loss function and their magnitude but also towards minimizing entropy. Once the training is complete, the weights will be arranged in such a way as to satisfy all three of these requirements.

However, at the end of the training, the weights will not be exactly in the quantized buckets, and each weight will need to be assigned to the corresponding bucket. This means that if, at the end of the training, the weight w_i lies within the interval $S_b = [\overline{w} - r + b\frac{2r}{C}, \overline{w} - r + (b+1)\frac{2r}{C})$ (or $S_b = [\frac{b}{C}, \frac{b+1}{C})$ in the particular simple case in which $\overline{w} = r = \frac{1}{2}$), it will be assigned the central value $\overline{w} - r + (2b+1)\frac{r}{C}$ (or $\frac{2b+1}{2C}$ in the aforementioned case), i.e., it will be quantized to that value.

This will further decrease the entropy without compromising the model's predictive capability.

Once this procedure is carried out, the network is ready to be compressed. It is, in fact, a sequence of weights (numerical values) for which there are numerous compression techniques to represent it with a smaller size in terms of occupied bytes. For this work, Huffman Coding and Arithmetic Coding were used. The latter algorithm, although much slower than the former, achieves entropy in the encoding, and it is the main reason why we have used this metric as a representative of the network size up to this point.

We also attempted to compress the network using more sophisticated compression algorithms such as gzip-9 and zstd-22, achieving even better results with zstd-22 compared to Huffman Coding and Arithmetic Coding. In fact, zstd-22 surpassed the 0-order entropy.

4 RESULTS

In this chapter, we present the results achieved by testing METaQ.

4.1 LeNet-5 & MNIST

The tests were conducted by training LeNet-5, a pioneering convolutional neural network (CNN) designed by Yann LeCun in 1998 [LeCun et al., 1998], primarily used for image recognition and, in particular, for classifying handwritten digits such as those in the MNIST dataset (short for Modified National Institute of Standards and Technology), a dataset widely used in the fields of machine learning and deep learning, especially for image classification.

MNIST is one of the most well-known datasets for evaluating optical character recognition (OCR) algorithms and contains images of handwritten digits representing the numbers 0 to 9. Each image is grayscale, with a size of 28x28 pixels (a total of 784 pixels), and is accompanied by a label indicating the represented digit (for example, an image of a "5" has the label 5). The training set consists of 60,000 images for training the models, while the test set contains 10,000 images to evaluate their performance.

LeNet-5 consists of a combination of convolutional layers and pooling (or subsampling) layers, followed by fully connected layers. Its architecture is relatively simple compared to modern CNNs, but it was revolutionary at the time of its creation. Here is a summary of its structure:

- Input Layer: Accepts grayscale images of size 32x32 (a larger format than MNIST, which is 28x28, so the images are usually resized). Each pixel of the image represents one unit in the input layer.

- Convolutional Layer 1: Applies 6 convolutional filters (kernels) of size 5x5. Output: 6 feature maps of size 28x28. The filters allow for the extraction of local features such as edges and textures.

- Pooling Layer 1: Applies a subsampling operation (average pooling) with 2x2 windows and a stride of 2. Output: 6 feature maps of size 14x14. Reduces the dimensionality and makes the features more robust to translations.

- Convolutional Layer 2 (C3): Applies 16 convolutional filters of size 5x5, with partial connections between feature maps (not all filters are connected to all input maps, reducing complexity). Output: 16 feature maps of size 10x10.

- Pooling Layer 2 (S4): Applies another layer of average pooling with 2x2 windows. Output: 16 feature maps of size 5x5.

- Fully Connected Layers: Layer 5: A fully connected layer with 120 neurons. Layer 6: Another fully connected layer with 84 neurons.

- Output Layer: Uses a softmax function to classify the image into 10 classes (from 0 to 9, for digit classification).

4.2 METaQ Test

Given the large number of hyperparameters in the METaQ strategy, a simple grid search in a desperate attempt to find a good configuration within the potentially vast range of possible values is certainly not a good idea. It is important to remember that training with the addition of the entropic term requires a significant amount of time ranging, depending on the hyperparameter C, from 30 seconds ($C = 6$) to 5 minutes per epoch ($C = 256$). This means that a single training run (with the computing resources available to us) could take several hours.

Therefore, to appropriately configure the hyperparameters, it is necessary to make an important assumption, which, while it may seem unrealistic at first, is actually quite reasonable and useful:

Many of these hyperparameters are independent of each other. This means that it is possible to tune one hyperparameter without worrying about the others. While this assumption may not hold perfectly for certain pairs of hyperparameters, for others, it is entirely reasonable.

In the following, we refer to Appendix 1 for a summary of the meaning of all the model's hyperparameters.

With reference to the problem (8), all the tests used:

$$L(w) = -\frac{1}{N} \sum_{j=1}^{N} \sum_{j'=1}^{M} y_{j,j'} \log(\hat{y}_{j,j'}(w)) \tag{88}$$

with $M = 10$ because in MNIST we have 10 classes, and

$$R(w) = ||w||_2 \tag{89}$$

for most of the cases[10]. That is, the cross-entropy was chosen as the loss function and the l_2 norm as the standard regularization term.

We mainly tested two optimizers for network training: ADAM and SGD. To tune the learning rate of these optimizers, we used the above independence assumption. Therefore the network was trained without the entropic regularization term, and the best value found was used in all the remaining tests. We found such a learning rate as $l_r = 0.0007$, which is actually a common value for this particular

---

[10]At the end of the tests we tried also to use the l_1 norm without having significant changes w.r.t. the $l2$ one.

network.

We did not perform tuning on the specific parameters of the Adam algorithm that control the exponential decay coefficients for the first-order moments (mean of the gradients) and second-order moments (variance of the gradients), by setting them to the standard values of 0.9 for the first and 0.999 for the second, which still led to good results.

To tune the parameters ζ_{FISTA} and ζ_{PBM} we did as follow: an instance of the dual problem was separately solved using the DSS stepsize method (which, as shown in Proposition 4, converges to the optimum) with a high number of iterations (1000). Now this optimum was used as a reference to compare it with the optimum found by the 3 algorithms (DSS stepsize, FISTA and PBM) with a reasonable number of iterations. We chose 20 as such a value, because a too large value could compromise performances once inserted into the training phase. In Tables 1-6 we can see those results. All of the results reported in Tables 1-6 are obtained with $\zeta_{FISTA} = 9e4$ and for $\zeta_{PBM} = 1e6$ which are just reasonable values (found after lots of attempts) but are not optimal in any strict sense. To optimize them it is necessary to tune them with respect to the particular value of C, therefore Tables 1-6 are just to show the impact of the 3 algorithms in the resolution of the dual lagrangian problem of maximization of $\phi_w(\xi)$.

	DSS stepsize (1000 iterations)	DSS stepsize (20 iterations)	FISTA	PBM
max $\phi_w(\xi)$	-5.47e7	-1.86e8	1.28e5	2.21e5
Execution Time [sec]	22.81	2.53	2.28	2.29

Table 1: C = 256, n = 44000

	DSS stepsize (1000 iterations)	DSS stepsize (20 iterations)	FISTA	PBM
max $\phi_w(\xi)$	-8.7e7	-2.14e8	2.41e5	2.12e4
Execution Time [sec]	13.19	1.23	1.39	1.32

Table 2: C = 128, n = 44000

	DSS stepsize (1000 iterations)	DSS stepsize (20 iterations)	FISTA	PBM
max $\phi_w(\xi)$	-6.00e5	-8.22e6	4.10e5	2.09e4
Execution Time [sec]	7.53	0.76	0.85	0.92

Table 3: C = 64, n = 44000

	DSS stepsize (1000 iterations)	DSS stepsize (20 iterations)	FISTA	PBM
max $\phi_w(\xi)$	-2.02e6	-6.32e7	4.52e5	2.21e4
Execution Time [sec]	5.72	0.56	0.57	0.57

Table 4: C = 32, n = 44000

31

	DSS stepsize (1000 iterations)	DSS stepsize (20 iterations)	FISTA	PBM
$\max \phi_w(\xi)$	-8.35e5	-1.86e7	4.99e5	2.64e5
Execution Time [sec]	4.31	0.45	0.46	0.45

Table 5: C = 16, n = 44000

	DSS stepsize (1000 iterations)	DSS stepsize (20 iterations)	FISTA	PBM
$\max \phi_w(\xi)$	-1.74e5	-7.92e6	5.46e5	2.05e4
Execution Time [sec]	3.30	0.36	0.35	0.36

Table 6: C = 8, n = 44000

To construct these tables we decided to perform the simulations with $n = 44000$ because the number of weights in the network we examined in these tests, LeNet-5, is approximately that value. As for C, we explored powers of 2 ranging from 8 to 256.

By analyzing these tables, we observe that as C decreases, the (absolute) value of the function's optimum tends to decrease, as do the execution times, which was expected. The execution times for algorithms with the same number of iterations (the last three columns) remain very similar, but what stands out is the optimum value found by the more sophisticated FISTA and PBM algorithms.

In fact, this value is significantly higher than what DSS stepsize manages to find. During training, each time a Lagrangian dual problem is solved, it starts from the optimal value found in the previous step. For this reason, after several iterations, DSS stepsize would also eventually reach the optimal values found by FISTA and PBM.

What is surprising, then, is that these two algorithms immediately achieve better values than DSS stepsize, reaching the optimal value (or at least a good approximation of it) much faster during the various training iterations.

The tuning of these parameters was performed for the value of C that was the main focus ($C = 6$), exploring several orders of magnitude. We report here the values of $\zeta_{FISTA} = \zeta_{PBM} = 1e5$, but without any claim that they are optimal, as the goal was simply to find good values that ensured fast convergence.

In terms of results, both FISTA and PBM successfully solve METaQ efficiently. To achieve the best model (for which we refer again to Appendix 1 for details), FISTA was used. For completeness, we report that a well-performing parameter configuration for PBM is max iteration $MI_{PBM} = 15$, $B_s = 10$, and $\zeta_{PBM} = 1e5$.

In equation (8), the coefficient of the standard regularization term is given by $\lambda\alpha$, while the coefficient of the entropy regularization term is given by $\lambda(1 - \alpha)$. To tune these two parameters,

we first studied the behavior of the network without entropy regularization to understand the most appropriate value for $\lambda\alpha$, and then we added a small perturbation and gradually adjusted its intensity.

Indeed, if we denote by $L_1 = \lambda\alpha$ and $L_2 = \lambda(1-\alpha)$, we find that once the value for L_1 is determined without perturbation, we can keep it fixed while varying L_2 to progressively increase the entropy regularization. α and λ are given by $\lambda = L_1 + L_2$ and $\alpha = L_1/\lambda$.

Regarding the number of epochs n_{epochs}, the following strategy was adopted: setting a very high value and terminating training if certain accuracy and entropy targets were not met within a given number of epochs.

If none of the stopping criteria[11] were met, there was a risk of being trapped in a good but suboptimal training process. For this reason, the number of epochs had to be set high enough to allow the algorithm to explore different configurations but not so high as to become stuck. A value of 100 was considered the minimum for a reasonable exploration, while training was never extended beyond 1000 epochs. Training the network for 1000 epochs with $C = 6$, the fastest configuration, takes approximately 500 minutes.

As for the parameter C, it is one of the most important parameters in the entire model but also one of the ones that most significantly impacts the execution time of the whole algorithm. As we already said, the execution time for an epoch with $C = 6$ is around 30 seconds, while for $C = 256$ it is around 300 seconds, namely an order of magnitude higher[12]. For this reason, we chose to keep it at low values ($C = 6$) with the idea of adjusting the actual number of quantization buckets in the post-training phase.

For the weight initialization, we also observed how the weights were distributed in the network without entropy regularization and tried to set the parameters $\overline{w}$ and r accordingly. These parameters are characteristic of each network and tell us within which range the weights should be initialized for optimal convergence. The LeNet-5 network for MNIST predictions without perturbation tends to converge its weights approximately in the range $[-1, 1]$ ($\overline{w} = 0, r = 1$), and it is around these values that we explored these two hyperparameters.

With the best parameters configuration (see Appendix 1 for all the details) by the 144th epoch, we achieved a model with 98.95% accuracy and an entropy of 114234 bits.

Quantization was performed by varying C within the range $[1, 1024]$, and the best results were those presented in Table 7. To differentiate the C used during training from the one used in the quantization phase, we will denote the latter as C^Q.

---

[11]A couple of examples of stopping criteria could be: If we are at the first epoch and the accuracy is below a certain starting threshold, stop. If we are at a certain epoch and the entropy is yet above a certain threshold, stop. One can also use more than one condition on the entropy and more than one condition on the accuracy.

[12]This execution times refer to simulations run on a spartan battle Macbook Pro with M1 chip.

$\mathbf{C}^Q$	Accuracy [%]	Entropy [bits]
140	99.01	59601
386	99.01	70135
331	99.01	68712
327	99.00	68583
322	99.00	68531
678	98.99	76202
220	98.99	63668
350	98.99	69694
377	98.99	70439

Table 7: Top-9 best results for the post-training quantization of the best model.

We see that for $C^Q = 140$, we achieved 99.01% of accuracy and an entropy of 59601 bits. This shows that not only was there no loss of accuracy, but there was actually an improvement, increasing it from 98.95% to 99.01%.

At this point, the compression was performed using various compression algorithms (see Table 8), achieving, as previously mentioned, the best result with zstd-22, which surpassed the 0-order entropy, reducing the size from 59601 bits to 48824 bits.

	Huffman Coding	Arithmetic Coding	gzip-9	zstd-22
Dimension [bits]	60503	59601	61336	48824
Compression Ratio [%]	4.26	4.19	4.31	3.43

Table 8: Sizes of the best model compressed with the various algorithms.

Considering that the initial entropy of the network weights starts at 1421632 bits, this means that the model manages to train the network while compressing it by a factor of $29\times$ (compression ratio of 3.43%) without losing predictive capability.

For comparison with the state of the art in compression strategies for LeNet-5 on MNIST, we report the following results just to mention a few:

- [Han et al., 2016], where they first prune the network by learning only the important connections, next they quantize the weights to enforce weight sharing in an iterative way, and finally they apply Huffman coding.

- [Jin et al., 2019], where they develop an accuracy-loss expected neural network compression framework, which involves network pruning, error bound assessment, optimization for error bound configuration, and compressed model generation.

We also have works with results inferior to those achieved with the METaQ strategy, such as:

- [Liu and Liu, 2018], where an improved scheme for pruning operations in compression methods is proposed by analizing first the distribution of network connection so as to determine the pruning threshold, and then using the pruning method to delete connections whose weights are less than the threshold.

- [Wang et al., 2018], where inspired by the tensor ring factorization, they use Tensor Ring Networks (TR-Nets) to significantly compress both the fully connected layers and the convolutional layers of deep networks.

- [Moya Rueda et al., 2017], where they present an efficient and robust approach for pruning entire neurons by exploiting maxout units for combining neurons into more complex convex functions and by making use of a local relevance measurement that ranks neurons according to their activation on the training set for pruning them, with a final parameter reduction comparison between neuron and weight pruning.

- [Wang et al., 2024], where they propose an energy-aware model compression method for various dataflow types in hardware architectures, EDCompress (EDC), which recasts the optimization process to a multistep problem and solves it by reinforcement learning algorithms.

Reference	Compression Ratio	Accuracy
[Liu and Liu, 2018]	9.47% (10.56×)	98.89%
[Wang et al., 2018]	9.09% (11×)	99.21%
[Moya Rueda et al., 2017]	8% (12.5×)	99.0%
[Wang et al., 2024]	3.85% (26×)	-
METaQ	**3.43% (29×)**	99.01%
[Han et al., 2016]	2.56% (39×)	99.26%
[Jin et al., 2019]	1.79% (56×)	99.7%

Table 9: METaQ with the State-of-the-art Compression Algorithm and their performances. Blank spaces stands for articles in which the exact value of the accuracy of the compressed model is not stated.

As can be seen, our results fall within a mid-to-high range, very close to the most cutting-edge algorithms.

We summarize in Table 9 the results of all the works cited.

The code for this entire work is available at the following link https://github.com/cardiaa/METaQ and can be easily accessed by anyone. In case of errors or feedback, please contact the author at andreacardia2008@gmail.com.

5 CONCLUSIONS

In this thesis, we developed and analyzed METaQ, an innovative algorithm designed to integrate entropy minimization into the training phase of a neural network, combining advanced optimization techniques and tailored approaches. METaQ was tested on LeNet-5 and MNIST, a relatively simple architecture and dataset, and the applicability of the algorithm to larger networks or complex datasets remains to be explored. It will be essential to test METaQ on more advanced network architectures, such as ResNet, Transformer, and Large Language Models, to evaluate its effectiveness in modern contexts and on more challenging datasets, such as ImageNet or CIFAR-10.

The presence of several hyperparameters (such as λ, α, or the number of buckets C) requires careful tuning to achieve optimal results. This sensitivity makes the algorithm subject to revision to adapt from one network to another. However, the description of how this fine-tuning was done to train LeNet-5 should serve as a useful guide in case the goal of the application changes. Given the large number of hyperparameters, some of these were not fully explored due to computational costs and available resources and time.

Although Lagrangian relaxation techniques and specialized algorithms provided an effective solution for computing the subgradient of $\phi(w)$, the approximations introduced could compromise accuracy in cases where the data structure or the network configuration do not perfectly align with the underlying assumptions.

A crucial area for development could be the search for more advanced optimization algorithms to solve the problem (17), or the integration of METaQ with other compression strategies, such as knowledge distillation or structured pruning, which could lead to even more computationally efficient models and better compression results.

Another critical point on which the entire analysis in this work was based is the form of entropy. One could consider the problem of minimizing higher-order entropy (instead of zero-order entropy), treating the network weights as not independent. A promising direction could be the introduction of compression metrics that exploit the dependency between weights, for example, using FM-index.

Studying alternative regularizations (like SPR) and evaluating other regularization functions that reduce entropy, exploring the possibility of models that more naturally balance compression and accuracy, could also bring some benefits.

The adoption of METaQ could have a significant impact on the efficiency of artificial intelligence models, especially in resource-constrained environments such as mobile or embedded devices. However, further studies are needed to address the identified limitations and fully explore its potential.

6 ACKNOWLEDGEMENTS

I would like to thank my Professors, who gave me the opportunity, through their ideas and insights, to be able to work on this project, which is as innovative and topical as it is stimulating and engaging, allowing me to interface with cutting-edge research work and technologies in a rapidly changing world.

But most of all, I thank my mom, who gave me the opportunity to graduate and build the person I am.

Appendix 1: METaQ hyperparameters.

We summarize in this Appendix all the hyperparameters called out in the text that are necessary for configuring the algorithm with a brief description of them, and their value used to obtain the model with the best results. To achieve these results, the optimizer used for the training was $ADAM$, the algorithm used to solve the lagrangian dual was $FISTA$, the loss function was the cross-entropy, and the standard regularization term was the l_2 norm.

Hyperparameter	Description	Best value
l_r	learning rate in the training optimizer	0.0007
β_1	moving average of the gradient in ADAM for training	0.9
β_2	moving average of the gradient square in ADAM for training	0.99
C	number of buckets during training	6
C^Q	number of buckets for the quantization	140
λ	coefficient of the total regularization in METaQ	0.0015
α	percentage of standard regularization versus total regularization	0.533
ζ_{FISTA}	parameter for adjusting the convergence pitch in FISTA	1e5
ζ_{PBM}	parameter for adjusting the convergence pitch in PBM	-
B_s	bundle size for PBM	-
MI_{FISTA}	maximum number of iteration in FISTA	15
MI_{PBM}	maximum number of iteration in PBM	-
$\overline{w}$	central value around which the weights are initialized	-0.11
r	radius of initialization of weights	1.114
$\overline{c}_b$	upper bound for c_b^*	1e-2
$\underline{c}_b$	lower bound for c_b^*	n
n_{epochs}	number of epochs used to train the network	144

Table 10: Hyperparameters of which METaQ is composed and their value for the best results. Since $FISTA$ was used to achieve these results, we left blank spaces in the fields related to the hyperparameters of the PBM algorithm. We refer to the discussion in Chapter 4.2 for recommended values for this algorithm.

References

Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. *Gpt-4 technical report*. arXiv preprint arXiv:2303.08774, 2023.

Amir Beck and Marc Teboulle. *A Fast Iterative Shrinkage-Thresholding Algorithm for Linear Inverse Problems*, volume 2, pages 183–202. SIAM J. IMAGING SCIENCES, 2009a. doi: 10.1137/080716542.

Amir Beck and Marc Teboulle. *A fast iterative shrinkage-thresholding algorithm for linear inverse problems*, volume 2, pages 183–202. SIAM journal on imaging sciences, 2009b.

Matteo Cacciola, Antonio Frangioni, and Andrea Lodi. *Structured pruning of neural networks for constraints learning*, volume 57, page 107194. Operations Research Letters, 2024. doi: https://doi.org/10.1016/j.orl.2024.107194.

Hongrong Cheng, Miao Zhang, and Javen Qinfeng Shi. *A survey on deep neural network pruning: Taxonomy, comparison, analysis, and recommendations*. IEEE, 2024.

Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. *Palm: Scaling language modeling with pathways*, volume 24, pages 1–113. Journal of Machine Learning Research, 2023.

cvxpy.org. Cvxpy api documentation. URL `https://www.cvxpy.org/api_reference/cvxpy.html`.

Giacomo D'Antonio and Antonio Frangioni. *Convergence Analysis of Deflected Conditional Approximate Subgradient Methods*, volume 20, pages 357–386. SIAM J. Optim., 2009. doi: 10.1137/080718814.

Ning Ding, Yujia Qin, Guang Yang, Fuchao Wei, Zonghan Yang, Yusheng Su, Shengding Hu, Yulin Chen, Chi-Min Chan, Weize Chen, et al. *Parameter-efficient fine-tuning of large-scale pre-trained language models*, volume 5, pages 220–235. Nature Publishing Group UK London, 2023.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. *The llama 3 herd of models*. arXiv preprint arXiv:2407.21783, 2024.

Moe Elbadawi, Hanxiang Li, Abdul W Basit, and Simon Gaisford. *The role of artificial intelligence in generating original scientific research*, volume 652, page 123741. International Journal of Pharmaceutics, Elsevier, 2024.

Maxim Enis and Mark Hopkins. *From LLM to NMT: Advancing Low-Resource Machine Translation with Claude*. arXiv preprint arXiv:2404.13813, 2024.

Salvatore Claudio Fanni, Maria Febi, Gayane Aghakhanyan, and Emanuele Neri. *Natural language processing*, pages 87–99. Springer, 2023.

Andrea Farruggia, Paolo Ferragina, Antonio Frangioni, and Rossano Venturini. *Bicriteria Data Compression*, volume 48, pages 1603–1642. SIAM Journal on Computing, 2019. doi: 10.1137/ 17M1121457.

Paolo Ferragina, Giovanni Manzini, Veli Mäkinen, and Gonzalo Navarro. *An alphabet-friendly FM-index*, pages 150–160. International Symposium on String Processing and Information Retrieval, 2004.

Antonio Frangioni. *About Lagrangian Methods in Integer Optimization*, volume 139, pages 163–193. Annals of Operations Research, 2005. doi: https://doi.org/10.1007/s10479-005-3447-9.

Antonio Frangioni. *Standard Bundle Methods: Untrusted Models and Duality*, pages 61–116. Springer International Publishing, Cham, 2020. doi: 10.1007/978-3-030-34910-3_3.

Antonio Frangioni and Enrico Gorgone. *A library for continuous convex separable quadratic knapsack problems*, volume 229, pages 37–40. European Journal of Operational Research, 2013. doi: https: //doi.org/10.1016/j.ejor.2013.02.038.

Elias Frantar and Dan Alistarh. *Sparsegpt: Massive language models can be accurately pruned in one-shot*, pages 10323–10337. International Conference on Machine Learning, 2023.

Zihao Fu, Haoran Yang, Anthony Man-Cho So, Wai Lam, Lidong Bing, and Nigel Collier. *On the effectiveness of parameter-efficient fine-tuning*, volume 37, pages 12799–12807. Proceedings of the AAAI conference on artificial intelligence, 2023.

Song Han, Huizi Mao, and William J. Dally. *Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding*. 2016.

Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The elements of statistical learning: data mining, inference, and prediction*. Springer, 2017.

Simon Haykin. *Neural networks and learning machines*. Pearson Education India, 2009.

Sian Jin, Sheng Di, Xin Liang, Jiannan Tian, Dingwen Tao, and Franck Cappello. *DeepSZ: A Novel Framework to Compress Deep Neural Networks by Using Error-Bounded Lossy Compression*, page 159–170. HPDC '19. Association for Computing Machinery, New York, NY, USA, 2019. ISBN 9781450366700. doi: 10.1145/3307681.3326608.

J. E. Kelley, Jr. *The Cutting-Plane Method for Solving Convex Programs*, volume 8, pages 703–712. Journal of the Society for Industrial and Applied Mathematics, 1960. doi: 10.1137/0108053.

Will Knight. *Nvidia's 3, 000'Personal AI Supercomputer'Will Let You Ditch the Data Center. Wired*, 2025.

Akshit Kurani, Pavan Doshi, Aarya Vakharia, and Manan Shah. *A comprehensive comparative study of artificial neural network (ANN) and support vector machines (SVM) on stock forecasting*, volume 10, pages 183–208. Annals of Data Science, Springer, 2023.

Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. *Gradient-based learning applied to document recognition*, volume 86, pages 2278–2324. Proceedings of the IEEE, 1998.

Xiaoya Li, Xiaofei Sun, Yuxian Meng, Junjun Liang, Fei Wu, and Jiwei Li. *Dice loss for data-imbalanced NLP tasks.* arXiv preprint arXiv:1911.02855, 2019.

Chongyang Liu and Qinrang Liu. *Improvement of pruning method for convolution neural network compression*, pages 57–60. Proceedings of the 2018 2nd International Conference on Deep Learning Technologies, 2018.

Haokun Liu, Derek Tam, Mohammed Muqeeth, Jay Mohta, Tenghao Huang, Mohit Bansal, and Colin A Raffel. *Few-shot parameter-efficient fine-tuning is better and cheaper than in-context learning*, volume 35, pages 1950–1965. Advances in Neural Information Processing Systems, 2022.

Zechun Liu, Barlas Oguz, Changsheng Zhao, Ernie Chang, Pierre Stock, Yashar Mehdad, Yangyang Shi, Raghuraman Krishnamoorthi, and Vikas Chandra. *Llm-qat: Data-free quantization aware training for large language models.* arXiv preprint arXiv:2305.17888, 2023.

Zhenhua Liu, Yunhe Wang, Kai Han, Wei Zhang, Siwei Ma, and Wen Gao. *Post-training quantization for vision transformer*, volume 34, pages 28092–28103. Advances in Neural Information Processing Systems, 2021.

Xinyin Ma, Gongfan Fang, and Xinchao Wang. *Llm-pruner: On the structural pruning of large language models*, volume 36, pages 21702–21720. Advances in neural information processing systems, 2023.

LE Melkumova and S Ya Shatskikh. *Comparing Ridge and LASSO estimators for data analysis*, volume 201, pages 746–755. Procedia engineering, Elsevier, 2017.

Tom M. Mitchell. *Machine learning*, volume 1. McGraw-hill New York, 1997.

Fernando Moya Rueda, Rene Grzeszick, and Gernot A Fink. *Neuron pruning for compressing deep networks using maxout architectures*, pages 177–188. German Conference on Pattern Recognition, 2017.

Yurii Nesterov. *A method for solving the convex programming problem with convergence rate O (1/k2)*, volume 269, page 543. 1983.

Massimo Pappalardo and Mauro Passacantando. *Ricerca operativa*. Didattica e ricerca. Manuali. Pisa University Press, 2012. ISBN 9788867411542. Casalini id: 2935851.

Md Iqbal Quraishi, J Pal Choudhury, and Mallika De. *Image recognition and processing using artificial neural network*. 2012 1st international conference on recent advances in information technology (RAIT), (pp. 95-100). IEEE., 2012.

Michael Santacroce, Zixin Wen, Yelong Shen, and Yuanzhi Li. *What matters in the structured pruning of generative language models?* arXiv preprint arXiv:2302.03773, 2023.

Anton Maximilian Schäfer and Hans Georg Zimmermann. *Recurrent neural networks are universal approximators*, pages 632–640. Artificial Neural Networks–ICANN 2006: 16th International Conference, Athens, Greece, September 10-14, 2006. Proceedings, Part I 16, 2006.

Aaquib Syed, Phillip Huang Guo, and Vijaykaarti Sundarapandiyan. *Prune and tune: Improving efficient pruning techniques for massive language models*. The First Tiny Papers Track at ICLR 2023, 2023.

Hongzhi Tong. *Functional linear regression with Huber loss*, volume 74. Journal of Complexity, 2023.

Tim Van Erven and Peter Harremos. *Rényi divergence and Kullback-Leibler divergence*, volume 60, pages 3797–3820. IEEE Transactions on Information Theory, 2014.

Qi Wang, Yue Ma, Kun Zhao, and Yingjie Tian. *A comprehensive survey of loss functions in machine learning*, pages 1–26. Annals of Data Science, 2020.

Wenqi Wang, Yifan Sun, Brian Eriksson, Wenlin Wang, and Vaneet Aggarwal. *Wide compression: Tensor ring nets*, pages 9329–9338. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2018.

Zhehui Wang, Tao Luo, Rick Siow Mong Goh, and Joey Tianyi Zhou. *EDCompress: Energy-Aware Model Compression for Dataflows*, volume 35, pages 208–220. IEEE Transactions on Neural Networks and Learning Systems, 2024. doi: 10.1109/TNNLS.2022.3172941.

Mingyang Zhang, Hao Chen, Chunhua Shen, Zhen Yang, Linlin Ou, Xinyi Yu, and Bohan Zhuang. *Loraprune: Pruning meets low-rank parameter-efficient fine-tuning*. arXiv preprint arXiv:2305.18403, 2023.

Xunyu Zhu, Jian Li, Yong Liu, Can Ma, and Weiping Wang. *A Survey on Model Compression for Large Language Models*, volume 12, pages 1556–1577. Transactions of the Association for Computational Linguistics, 2024. doi: 10.1162/tacl_a_00704.